

Model Aggregation: Data-driven combination of black box models


Theo Bourdais

PhD Student, Computing + Mathematical Sciences

June 15th 2025



Real-life example from the IPCC



Arctic sea ice extent estimated by many models,


Coupled Model Intercomparison Project (report AR5 - figure 9.24)

The aggregation problem

Given the models,

create a combination
(i.e. aggregation)

that approximates the target



Best Mean Squared Error Aggregation

The best possible aggregation in Mean Squared Error is

$$M_A^*(x) := \underset{f \text{ measurable}}{\operatorname{argmin}} \mathbb{E}[|Y(x) - f(x, M_1(x), \dots, M_n(x))|^2] = \mathbb{E}[Y(x)|M_1(x), \dots, M_n(x)]$$

This is intractable in general

Special Case: $(Y(x), M_1(x), \dots, M_n(x))$ is Gaussian

$$M_A^*(x) = \sum_{i=1}^n \alpha_i^*(x) M_i(x)$$

$$\alpha^*(x) = \underset{a \in \mathbb{R}^n}{\operatorname{argmin}} \mathbb{E} \left[\left| Y(x) - \sum_{i=1}^n a_i M_i(x) \right|^2 \right] = \mathbb{E} [M(x) M(x)^T]^{-1} \mathbb{E} [M(x) Y(x)]$$

Best case aggregation: Gaussian models

To solve the Laplace equation:

$$\begin{cases} \Delta Y = f & \text{on } \Omega \\ Y = g & \text{on } \partial\Omega \end{cases}$$

We can use a Gaussian process with:

- A kernel k
- A set of collocation points $X \subset \Omega$


To get a **Gaussian** approximation of the solution [Chen et al., 2021]

$$\xi \sim \mathcal{N}(0, k) \quad \hat{Y} = \mathbb{E}[\xi \mid \Delta\xi(X) = f(X)]$$


Predictions

$$\text{Laplace equation} \begin{cases} \Delta Y = f & \text{on } \Omega \\ Y = g & \text{on } \partial\Omega \end{cases}$$

$M_1(x)$




$M_2(x)$




⋮


$M_n(x)$



Average



True solution




$Y(x)$

Caltech


Predictions

Laplace equation $\begin{cases} \Delta Y = f & \text{on } \Omega \\ Y = g & \text{on } \partial\Omega \end{cases}$

$M_1(x)$




$M_2(x)$




⋮

$M_n(x)$




$$\sum_{i=1}^n \alpha^*(x) M_i(x)$$



\approx

True solution



$Y(x)$

Caltech

Minimal Error Aggregation

This does not work!

α^* is defined as:

$$\alpha^*(x) = \underset{a \in \mathbb{R}^n}{\operatorname{argmin}} \mathbb{E} \left[\left| Y(x) - \sum_{i=1}^n a_i M_i(x) \right|^2 \right]$$

And we only have access to data $\{X_i, Y(X_i)\}_{i=1}^N$. So we could pick a Machine Learning Method, learn over the training set and extrapolate for all x

$$\hat{\alpha}_E = \underset{a}{\operatorname{argmin}} \sum_{k=1}^N \left[\left| Y(X_k) - \sum_{i=1}^n a_i(X_k) M_i(X_k) \right|^2 \right]$$

(This is Mixture-of-Experts with frozen experts)


A pathological example

Given the target, models and data:

- Take α linear
 $\alpha(x) = (a_Gx + b_G, a_Bx + b_B)$
- Train using empirical MSE

Notice that:

- For each data point, the good model performs better than the bad model



- The aggregate ignores the good model and interpolates the data
- Aggregation uses models as features, not approximations of Y

Minimal Variance Aggregation

Problem: we don't have enough constraints / we didn't define what a good model is.
Let

$$\begin{cases} M_1(x) = Y(x) + \epsilon_1(x) \\ \vdots \\ M_n(x) = Y(x) + \epsilon_n(x) \end{cases} \quad \text{where} \quad \begin{cases} \text{(For simplicity)} & \epsilon_i \text{ are independent} \\ \text{(Write)} & \text{Var}[\epsilon_i(x)] = V_i(x) \\ \text{(Assumption)} & \mathbb{E}[\epsilon_i(x)] = 0 \end{cases}$$

Then the aggregation is unbiased if

$$\sum_{i=1}^n \alpha_i = 1$$

Minimal Variance Aggregation

$$\alpha_V(x) = \operatorname{argmin} \mathbb{E} \left[\left| Y(x) - \sum_{i=1}^n \alpha_i M_i(x) \right|^2 \right] \quad \xrightarrow{\sum_{i=1}^n \alpha_i = 1} \quad \alpha_V(x)^T M(x) = \frac{\sum_{i=1}^n \frac{1}{V_i(x)} M_i(x)}{\sum_{i=1}^n \frac{1}{V_i(x)}}$$

Note that:

- If $\epsilon_i(x) \sim \mathcal{N}(0, V_i(x))$, this is MLE
- If $\mathbb{E}[\epsilon_i(x)] = 0$, $\mathbb{E}[\epsilon_i(x)^2] = V_i(x)$, this is BLUE
- If no assumption, best convex combination

We just need to learn $V_i(x)$, the expected error of each model

Learning the variance/error

To predict the variance, we:

- Write $V_i(x) = e^{\lambda_i(x)}$ where λ_i is a Machine Learning method (Gaussian process, neural network...) to ensure positivity
- Then the aggregation is a softmax
- Use the loss

$$\min_{\lambda_i \in \mathcal{H}} \sum_{k=1}^N \left[e^{\lambda_i(X_k)} - (Y(X_k) - M_i(X_k))^2 \right]^2 + \eta \|\lambda_i\|_{\mathcal{H}}^2$$

$V_i(X_k)$ Empirical variance Regularization

Theorem on linear regression:

Assume samples $(M_j, Y_j)_{j=1}^N$, which one has the best loss $\mathcal{L}(\alpha) = \mathbb{E}[|Y - \alpha^T M|^2]$?

$$\hat{\alpha}_E(x) = \operatorname{argmin}_{a \in \mathbb{R}^n} \sum_{j=1}^N \left[|Y_j - a^T M_j|^2 \right]$$

Minimal (Empirical) **Error** Aggregation

$$\hat{\alpha}_V(x) = \operatorname{argmin}_{a \in \mathbb{R}^n} \begin{cases} \sum_{j=1}^N \left[|Y_j - a^T M_j|^2 \right] \\ \text{such that } \sum_{i=1}^n a_i = 1 \end{cases}$$

Minimal (Empirical) **Variance** Aggregation

There exists $\lambda \in [0,1]$ s.t.:

$$\mathcal{L}(\hat{\alpha}_E) = \mathcal{L}(\alpha^*) + \mathcal{O}\left(\frac{1}{\sqrt{N}}\right)$$


$$\mathcal{L}(\hat{\alpha}_V) = \frac{1}{\lambda} \mathcal{L}(\alpha^*) + \mathcal{O}\left(\frac{1}{N}\right)$$

In model aggregation, N is small and $\lambda \rightarrow 1$

Applications

The Boston housing dataset


- Data: 506 samples $\{X_i, Y_i\}$
- Data is split into train-test-val
- Aggregation of **red models** using val data
 - Red models only see train data
 - Blue models for comparison see train+val
- Aggregation is:
 - Better than models aggregated
 - Better than the mean
 - Better than all models



The Boston housing dataset

A comparison with minimal error aggregation:

- Take two identical Neural networks
- Train:
 - To minimize error (bad loss)
 - To estimate variance (our loss)




PDE examples

Given a PDE, we may have multiple solvers/approximations giving a solution.


For example:

Laplace equation:
$$\begin{cases} \Delta u = f & \text{on } \Omega \\ u = 0 & \text{on } \partial\Omega \end{cases}$$

Given models $M_i(f) \approx u$, we want to learn the aggregation operator $\alpha(f)$




Random f




Random u

PDE example 1 - Laplace equation



GP diverged



Method	Geometric mean of MSE (log scale)
Aggregate	-6.282
FDM	-5.523
Spectral	-4.988
Gaussian process	-4.739
FDM asymmetric (right)	-4.685
FDM asymmetric (left)	-4.699

PDE example 2 - Burger's equation


Consider Burger's equation on $\Omega = [0,1]^2$:

$$\begin{cases} \partial_t u + u \partial_x u = \nu \partial_{xx} u & \text{for } (x,t) \in \Omega \\ u(0,x) = f(x) & \text{for } x \in [0,1] \\ u(t,0) = u(t,1) & \text{for } t \in [0,1] \end{cases}$$


Choose:

- ν to be small
- $f \sim \mathcal{N}(0,K)$ where $K(x,y) = \exp\left(-\frac{2}{l^2} \sin^2\left(\pi|x_i - x_j|^2\right)\right)$
 - i.e. f is periodic and infinitely differentiable


PDE example 2 - Burger's equation



Correct




Blowup




Oscillations

PDE example 2 - Burger's equation



Method	Geometric mean of MSE (log scale)
Aggregate	-3.106
Riemann	-2.734
TVD	-2.568
FDM	-1.228
Spectral	-0.625
Implicit	-0.488
Explicit	-0.455
Lax-Wendroff	-0.455



Conclusion

We introduce a simple framework to aggregate existing models

- Only requires model output (no assumption, non intrusive)
- Most useful in scientific computing settings with legacy models
- Aggregate any type of methods (ML, solvers...)

Bourdais, T., & Owhadi, H. (2025).


Minimal Variance Model Aggregation: A principled, non-intrusive, and versatile integration of black box models

ICLR 2025




Predictions

$M_1(x)$




$M_2(x)$




⋮

$M_n(x)$




We want to create this !

Model
Aggregation



≈

True solution



$Y(x)$

Caltech

Summary

Given $\begin{matrix} M_1(x) \\ \vdots \\ M_n(x) \end{matrix} \rightarrow f(x, M_1(x), \dots, M_n(x)) \approx Y(x)$

Summary

Given $M_1(x)$ \vdots $M_n(x)$ $\sum_{i=1}^n \alpha_i(x)M_i(x)$ Where $\approx Y(x)$

1. Simplification + Gaussian ideal case

$$\alpha^*(x) = \underset{a \in \mathbb{R}^n}{\operatorname{argmin}} \mathbb{E} \left[\left| Y(x) - \sum_{i=1}^n a_i M_i(x) \right|^2 \right]$$

Summary

Given $M_1(x) \quad \vdots \quad M_n(x)$

$$\sum_{i=1}^n \alpha_i(x) M_i(x)$$

- Where
1. Simplification + Gaussian ideal case
 2. Directly minimize error
Does not work

No assumption

$$\alpha^* \in \arg\min_a \sum_{i=1}^N \left[\left\| Y(X_i) - \sum_{i=1}^n \alpha_i M_i(X_i) \right\|^2 \right]$$



$$\{X_i, Y(X_i)\}_{i=1}^N$$

Summary

Given $M_1(x)$ \vdots $M_n(x)$

$$\sum_{i=1}^n \alpha_i(x) M_i(x) \approx Y(x)$$

$\mathbb{E}[M_i(x)] = 0$

Where $\alpha_i(x) = \frac{1}{Var[M_i(x)]} \sum_{k=1}^n \frac{1}{Var[M_k(x)]}$

1. Simplification + Gaussian ideal case
2. Directly minimize error
Does not work
3. Assume unbiased models

Summary

Given

$$\begin{matrix} M_1(x) \\ \vdots \\ M_n(x) \end{matrix} \quad \sum_{i=1}^n \alpha_i(x) M_i(x)$$

$$\mathbb{E}[M_i(x)] = 0$$

Where $\alpha_i(x) = \frac{\frac{1}{Var[M_i(x)]}}{\sum_{k=1}^n \frac{e^{-\lambda_k(x)}}{Var[M_k(x)]}}$

$$\lambda_i = \operatorname{argmin}_{l \in \mathcal{H}} \sum_{k=1}^N \left[e^{l(X_k)} - (Y(X_k) - M_i(X_k))^2 \right]^2 + \eta \|l\|_{\mathcal{H}}^2$$



ML regression

(Neural network, Gaussian process...)

$$\{X_i, Y(X_i)\}_{i=1}^N$$

1. Simplification + Gaussian ideal case
2. Directly minimize error
Does not work
3. Assume unbiased models
4. Learn
 $e^{\lambda_i(x)} \approx Var[M_i(x)]$

Minimal Variance Aggregation

Let:


$$\begin{cases} M_1(x) = Y(x) + \epsilon_1(x) \\ \vdots \\ M_n(x) = Y(x) + \epsilon_n(x) \end{cases}$$

Where


- ϵ_i are independent (ease of presentation)
- We write $\mathbb{E} [|Y(x) - M_i(x)|^2] = V_i(x)$

Predictions

$M_1(x)$




$M_2(x)$




⋮


$M_n(x)$



Average



True solution



$Y(x)$

Caltech